

ENTRUST: Engineering Trustworthy Self-Managing Systems – Technical Report

1 INTRODUCTION

Since the advent of autonomous computing and the vision of autonomous systems the engineering of self-managing software systems has been a research challenge. Self-managing systems are expected not only to operate with minimal administrator intervention throughout their lifetime, but also to adapt while providing service in response to environment changes, evolving requirements and unexpected failures. As a result, a lot of research has been devoted in developing frameworks, methodologies and architectures that support the engineering of systems enhanced with self-* capabilities(e.g., self-adaptive, self-managing, self-healing).

Simos References
needed

Despite the research progress achieved within the last few years, the engineering of trustworthy self-managing systems remains a major research challenge. This barrier constrains the applicability of self-managing systems in safety-critical and business-critical applications, as for example, in healthcare, e-commerce, and defence. Systems deployed in these application areas should work dependably and must be characterised by high-integrity runtime operation as defined by NIST: “High integrity software is software that must be trusted to work dependably in some critical function, and whose failure to do so may have catastrophic results, such as serious injury, loss of life or property, business failure or breach of security.”

The software engineering community classified the provision of assurances that a system operates dependably among the most important research objectives for self-managing systems. In fact, assurances was amid the research threads highlighted in the most recent roadmaps for self-adaptive systems. We adopt the following definition for assurances: “*Assurances is providing evidence that the software system fulfils its functional and non-functional requirements throughout its lifetime*”.

In this work, we extend the current practices on the provision of assurances for self-managing systems. To this end, we introduce the first tool-supported methodology for assuring the safety of the closed control loops of self-managing systems using established industry practices. Our framework for the ENgineering of TRUstworthy Self-managing sysTems (ENTRUST) integrates:

- an extended version of our approach to developing formally verified monitor-analysis-planning-execution (MAPE) autonomic control loops;
- our runtime quantitative verification technique for the analysis stage of these loops;
- an industry-adopted approach to arguing security based on the widely used Goal Structuring Notation (GSN) for safety arguments.

2 Running Example

We will use a self-adaptive UUV system to evaluate ENTRUST. UUVs are increasingly used in a wide range of oceanographic and military tasks, including oceanic surveillance (e.g., to monitor pollution levels and ecosystems), undersea mapping, and mine detection. Limitations intrinsic to the environment in which these vehicles operate (e.g., impossibility to maintain UUV-operator communication during missions and high frequency of unexpected changes) require that UUV systems are self-adaptive. These systems are also safety critical (e.g., when used for mine detection and surveillance of ecosystems that should not be impacted) and/or business critical, since UUVs are often expensive equipment that should not be lost during missions.

The self-adaptive system in our study consists of a UUV used to carry out a surveillance and data gathering mission. The UUV is equipped with $n \geq 1$ on-board sensors that can measure the same attribute of the ocean environment (e.g., water current, salinity or thermocline). When used, the sensors take measurements with different, variable rates r_1, r_2, \dots, r_n . The probability that each sensor produces measurements that are sufficiently accurate for the purpose of the mission depends on the UUV speed sp , given by p_1, p_2, \dots, p_n ¹. For each measurement taken, different amount of energy is consumed, given by e_1, e_2, \dots, e_n . Finally, the n sensors can be switched on and off individually (e.g., to save battery power when not required), but these operations consume an amount of energy given by $e_1^{\text{on}}, e_2^{\text{on}}, \dots, e_n^{\text{on}}$ and $e_1^{\text{off}}, e_2^{\text{off}}, \dots, e_n^{\text{off}}$, respectively.

The UUV is required to self-adapt to changes in the observed sensor measurement rates r_i , $1 \leq i \leq n$, and to sensor failures by dynamically adjusting:

- (a) the UUV speed sp
- (b) the sensor configuration x_1, x_2, \dots, x_n (where $x_i = 1$ if the i -th sensor is on and $x_i = 0$ otherwise)

so that the UUV complies with the following requirements at all times:

- R1:** The UUV should take at least 20 measurements of sufficient accuracy for every 10 metres of mission distance.
- R2:** The energy consumption of the sensors should not exceed 120 Joules per 10 surveyed metres.
- R3:** If requirements R1 and R2 are satisfied by multiple configurations, the UUV should use one of these configurations that minimises the cost function

$$cost = w_1 E + w_2 s^{-1}, \quad (1)$$

where E represents the energy consumed by the sensors to survey a 10m mission distance, and $w_1, w_2 > 0$ represent weights that reflect the relative importance of carrying out the mission with reduced battery usage and completing the mission faster.

¹ This information can be extracted from the technical specification of sensors; for example, see <http://www.ashtead-technology.com/rental-equipment/teledyne-rdi-600khz-navigator/>

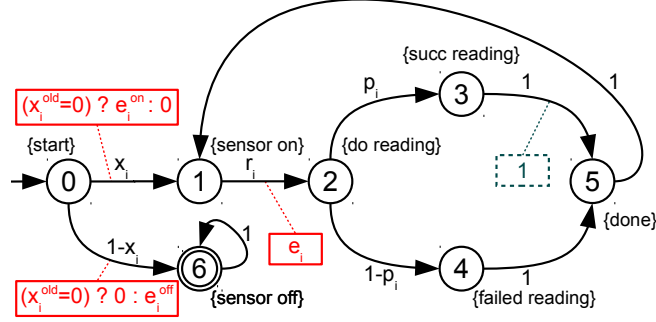


Fig. 1. CTMC model M_i of the i -th UUV sensor

Example 1. Figure 1 depicts the CTMC model M_i of the i -th sensor of the UUV from our running example. The CTMC starts in state 0 and transitions in state 1 if the sensor is switched on ($x_i = 1$) or in state 6 otherwise ($x_i = 0$). With rate r_i , a measurement is performed, as indicated by the transition between states 1 and 2. The measurement is “sufficiently accurate” with probability p_i and the transition between states 2 and 3 is taken, whereas the transition between states 2 and 4 is enabled in case of an inaccurate measurement. An active sensor continues taking measurements while it is active, as modelled by the transition between states 5 and 1. The CTMC model is augmented with two cost/reward structures, whose non-zero elements are shown in Fig. 1 in rectangular boxes, and dashed rectangular boxes, respectively. The former, “*energy*” structure associates the energy used to switch the sensor on (i.e., e_i^{on}) and off (i.e., e_i^{off}) and to perform a measurement (i.e., e_i) with the CTMC transitions that model these events. Note that the ternary expression “*condition*?*a*:*b*” was used to indicate that the energy e_i^{on} is consumed only if the previous state of the sensor was “off”, i.e., if $x_i^{old} = 0$, whereas the energy e_i^{off} is used only if the opposite is true ($x_i^{old} = 1$). As concerns the latter, “*measurement*” cost/reward structure, it associates a reward of 1 with the transition that corresponds to an accurate measurement.

3 ActivFORMS

ActivForms

4 Runtime Quantitative Verification

Runtime quantitative verification (RQV) [?] is an approach to implementing the closed control loop of self-adaptive systems in which the adaptation decisions are driven by the continual verification of stochastic models. RQV was introduced in [?, ?] and further refined in [?, ?, ?].

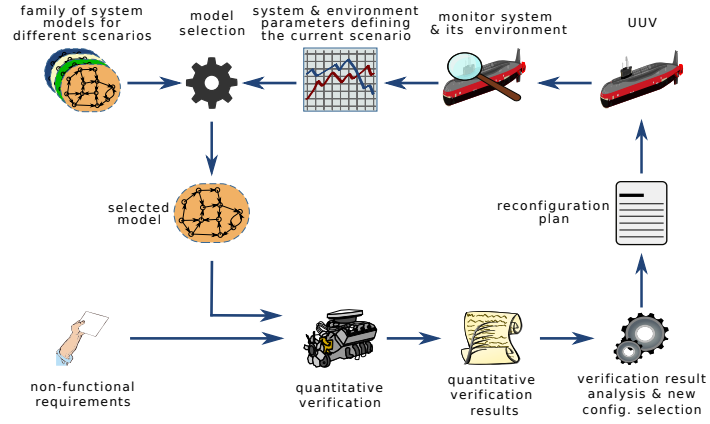


Fig. 2. RQV-based self-adaptive system

Figure 2 depicts the adaptation workflow of an RQV-based self-adaptive system. The approach involves monitoring the system (e.g., the UUV) and its environment continuously, to identify relevant changes and quantify them using fast on-line learning techniques. When these changes are significant and/or periodically, an RQV step is triggered, to identify which scenario the system operates in, and to select the associated model from a family of system models (i.e., a *parametric model*) that correspond to different such scenarios. The selected model is then analysed using quantitative verification, to identify and/or predict violations of QoS requirements such as response time, availability and cost. When QoS violations are identified or predicted, the verification results support the selection of a new set of values for the configurable system parameters, so that enforcing this new configuration is *guaranteed* to restore or maintain compliance with QoS requirements, respectively.

5 Implementation

This section will explain how we achieved our goal.

5.1 Monitor

Model of Monitor template

5.2 Analyzer

Model of Analysis template

5.3 Planner

Model of Planning template

5.4 Executor

Model of Execution template

5.5 Probes & Effectors

If necessary code of probes & effectors

5.6 Offline Assurances

The provision of partial offline assurances

5.7 Online Assurances

The perpetual provision of assurances online. How we are using both technologies together.

6 Methodology